

A stable learner is a prediction model that generalises to new samples without re-evaluation of the current kernel. In market settings, a learner is stable if, and only if, its pricing kernel is arbitrage-free. Stable pricing kernels preserve information and support the efficient markets hypothesis since the marginal impact of the next observation is decreasing in the relative frequency of its type. Our contribution is to derive the current kernel on the basis of sample-based sentiments (rankings) for a broader class of stable learner. This extension of Gilboa and Schmeidler (2003) is justified by US Treasury bond market data for the period 1961-2023. Our “partial-3-diversity of data” condition on daily rankings of annual yields holds across 91% of the sample for all 30 maturity dates. When stronger diversity conditions fail, stability requires out-of-sample validation: sentient agents might simulate or imagine the impact of novel data; artificial ones might engage in leave-one-type-out cross validation; or, one might bypass sentiments and directly rule out Dutch books. Our framework embeds current sentiments in a system of potential generalisations thus enabling within-sample testing of stability for any form of external validation.